

Student: Balakrishna Mupparaju

Assignment: Week5 & 6

Activity 5.01: Reading Tabular Data from a Web Page and Creating DataFrames

```
In [4]: # Import necessary libraries
from bs4 import BeautifulSoup # To parse HTML files
import pandas as pd # For managing tabular data and creating DataFrames

# Open and read the local HTML file containing GDP data
fd = open("/Users/balakrishnamupparaju/Downloads/The-Data-Wrangling-Workshop.html")
soup = BeautifulSoup(fd) # Parse the file's contents using BeautifulSoup
fd.close() # Close the file to release system resources

# Find all tables within the HTML document
all_tables = soup.find_all("table")
print("Total number of tables are {}".format(len(all_tables))) # Print total number of tables
```

Total number of tables are 9

```
In [16]: # Select a specific table based on its class
data_table = soup.find("table", {"class": "wikitable"})
print("Type of the selected table is: {}".format(type(data_table))) # Confirm the type of the selected table

# Extract the first row (likely header-like or meta-data row) of the table
sources = data_table.tbody.findAll('tr', recursive=False)[0]

# Extract all <td> elements (columns) from the first row
sources_list = [td for td in sources.findAll('td')]
print("Number of source columns found: {}".format(len(sources_list))) # Print the number of source columns found

# Extract the second row (likely the first row of data) in the table
data = data_table.tbody.findAll('tr', recursive=False)[1].findAll('td', recursive=False)

# Initialize an empty list to store nested tables found within the main table
data_tables = []
for td in data:
    data_tables.append(td.findAll('table')) # Find any nested tables in each row
print("Number of nested tables found: {}".format(len(data_tables))) # Print the number of nested tables found

# Extract source names from the header row by finding anchor tags (<a>)
source_names = [source.findAll('a')[0].getText() for source in sources_list]
print("Source names extracted: {}".format(source_names)) # Print the source names
```

Type of the selected table is: <class 'bs4.element.Tag'>
Number of source columns found: 3
Number of nested tables found: 3
Source names extracted: ['International Monetary Fund', 'World Bank', 'United Nations']

```
In [17]: # ----- FIRST TABLE -----
# Extract headers (column names) from the first nested table
header1 = [th.getText().strip() for th in data_tables[0][0].findAll('thead')]
print("Headers for the first table: {}".format(header1)) # Print the header

# Extract all rows of data from the first nested table
rows1 = data_tables[0][0].findAll('tbody')[0].findAll('tr')
data_rows1 = [[td.getText().strip() for td in tr.findAll('td')] for tr in rows1]

# Create a DataFrame for the first table
df1 = pd.DataFrame(data_rows1, columns=header1)
print("First DataFrame preview:")
print(df1.head()) # Display the first few rows of the DataFrame
```

Headers for the first table: ['Rank', 'Country', 'GDP(US\$MM)']

First DataFrame preview:

	Rank	Country	GDP(US\$MM)
0		World[19]	79,865,481
1	1	United States	19,390,600
2	2	China[n 1]	12,014,610
3	3	Japan	4,872,135
4	4	Germany	3,684,816

```
In [18]: # ----- SECOND TABLE -----
# Extract headers from the second nested table
header2 = [th.getText().strip() for th in data_tables[1][0].findAll('thead')]
print("Headers for the second table: {}".format(header2)) # Print the header

# Extract all rows of data from the second nested table
rows2 = data_tables[1][0].findAll('tbody')[0].findAll('tr')

# Define a helper function to clean and extract the desired text
def find_right_text(i, td):
    if i == 0: # First column logic
        return td.getText().strip()
    elif i == 1: # Second column logic
        return td.getText().strip()
    else: # For other columns, extract text after the ♠ symbol
        index = td.text.find("♠")
        return td.text[index+1:].strip()

# Apply the helper function to process all data rows
data_rows2 = [[find_right_text(i, td) for i, td in enumerate(tr.findAll('td'))] for tr in rows2]

# Create a DataFrame for the second table
df2 = pd.DataFrame(data_rows2, columns=header2)
print("Second DataFrame preview:")
print(df2.head()) # Display the first few rows of the DataFrame
```

Headers for the second table: ['Rank', 'Country', 'GDP(US\$MM)']

Second DataFrame preview:

	Rank	Country	GDP(US\$MM)
0		World	80,683,787
1	1	United States	19,390,604
2		European Union[23]	17,277,698
3	2	China[n 4]	12,237,700
4	3	Japan	4,872,137

```
In [19]: # ----- THIRD TABLE -----
# Extract headers from the third nested table
header3 = [th.getText().strip() for th in data_tables[2][0].findAll('thead')]
print("Headers for the third table: {}".format(header3)) # Print the header

# Extract all rows of data from the third nested table
rows3 = data_tables[2][0].findAll('tbody')[0].findAll('tr')

# Apply the helper function to process all data rows
data_rows3 = [[find_right_text(i, td) for i, td in enumerate(tr.findAll('td'))] for tr in rows3]

# Create a DataFrame for the third table
df3 = pd.DataFrame(data_rows3, columns=header3)
print("Third DataFrame preview:")
print(df3.head()) # Display the first few rows of the DataFrame
```

Headers for the third table: ['Rank', 'Country', 'GDP(US\$MM)']

Third DataFrame preview:

	Rank	Country	GDP(US\$MM)
0		World[24]	75,648,448
1	1	United States	18,624,475
2	2	China[n 4]	11,218,281
3	3	Japan	4,936,211
4	4	Germany	3,477,796

Activity 6.01: Handling Outliers and Missing Data

In this activity, we will identify and get rid of outliers. Here, we have a CSV file. The goal here is to clean the data by using the knowledge that we have learned about so far and come up with a nicely formatted DataFrame. Identify the type of outliers and their effect on the data and clean the messy data.

```
In [28]: # Import required libraries
import pandas as pd # For working with tabular data
import numpy as np # For numerical operations
import matplotlib.pyplot as plt # For visualizing data

# Enable inline plotting for visualizing graphs directly in notebooks
%matplotlib inline

# Step 1: Load the dataset from the given CSV file
df = pd.read_csv("/Users/balakrishnamupparaju/Downloads/The-Data-Wrangling-w")
```

```
print("Loaded dataset preview:")
print(df.head()) # Display the first 5 rows of the dataset
```

Loaded dataset preview:

	id	first_name	last_name	email	gender	\
0	1	Sonny	Dahl	sdahl0@mysql.com	Male	
1	2	NaN	NaN	dhoovart1@hud.gov	NaN	
2	3	Gar	Armal	garmal2@technorati.com	NaN	
3	4	Chiarra	Nulty	cnulty3@newyorker.com	NaN	
4	5	NaN	NaN	sleaver4@elegantthemes.com	NaN	

	ip_address	visit
0	135.36.96.183	1225.0
1	237.165.194.143	919.0
2	166.43.137.224	271.0
3	139.98.137.108	1002.0
4	46.117.117.27	2434.0

```
In [29]: # Step 2: Check for duplicate entries in important columns
print("First name has duplicates: {}".format(any(df.first_name.duplicated())))
print("Last name has duplicates: {}".format(any(df.last_name.duplicated())))
print("Email has duplicates: {}".format(any(df.email.duplicated())))
```

First name has duplicates: True

Last name has duplicates: True

Email has duplicates: False

```
In [30]: # Step 3: Check for missing (NaN) values in critical columns
print("Email column contains missing values: {}".format(df.email.isnull().value_counts()))
print("IP Address column contains missing values: {}".format(df.ip_address.isnull().value_counts()))
print("Visit column contains missing values: {}".format(df.visit.isnull().value_counts()))
```

Email column contains missing values: False

IP Address column contains missing values: False

Visit column contains missing values: True

```
In [32]: # Step 4: Record the size of the original dataset for comparison
size_prev = df.shape # Get the shape of the DataFrame (rows, columns)

# Step 5: Remove rows where the 'visit' column has NaN or non-numeric values
df = df[np.isfinite(df['visit'])] # Filter rows with finite numeric values
size_after = df.shape # Get the new shape after cleaning
```


```
In [33]: # Step 6: Print size comparison of the DataFrame before and after cleaning
print("Original data size: {} rows.".format(size_prev[0]))
print("Data size after removing NaN/invalid 'visit' values: {} rows.".format(size_after[0]))
```

Original data size: 1000 rows.

Data size after removing NaN/invalid 'visit' values: 974 rows.

```
In [34]: # Step 7: Visualize the 'visit' column to identify potential outliers
plt.boxplot(df.visit, notch=True)
plt.title("Boxplot of 'visit' column")
plt.show()
```

#Boxplot Visualization: Helps in identifying the distribution and spotting outliers



```
In [35]: # Step 8: Remove outliers based on 'visit' value thresholds
# Keep rows where 'visit' values are within a reasonable range (100 to 2900)
df1 = df[(df['visit'] <= 2900) & (df['visit'] >= 100)]
print("Data size after removing outliers: {} rows.".format(df1.shape[0]))
```

Data size after removing outliers: 923 rows.

```
In [36]: # Final output preview
print("Cleaned dataset preview:")
print(df1.head())
```

Cleaned dataset preview:

	id	first_name	last_name	email	gender	\
0	1	Sonny	Dahl	sdahl0@mysql.com	Male	
1	2	NaN	NaN	dhoovart1@hud.gov	NaN	
2	3	Gar	Armal	garmal2@technorati.com	NaN	
3	4	Chiarra	Nulty	cnulty3@newyorker.com	NaN	
4	5	NaN	NaN	sleaver4@elegantthemes.com	NaN	

	ip_address	visit
0	135.36.96.183	1225.0
1	237.165.194.143	919.0
2	166.43.137.224	271.0
3	139.98.137.108	1002.0
4	46.117.117.27	2434.0

In []:

3. Insert data into a SQL Lite database

```
In [2]: # contacts_sqlite.py

import sqlite3
import pandas as pd

# -----
# Step 1: Connect to SQLite (in-memory or file-based)
# -----
conn = sqlite3.connect("contacts.db") # Saves as contacts.db file
cursor = conn.cursor()

# -----
# Step 2: Create the contacts table
# -----
cursor.execute("""
CREATE TABLE IF NOT EXISTS contacts (
    name TEXT,
    address TEXT,
    city TEXT,
    state TEXT,
    zip TEXT,
    phone_number TEXT
)
""")

# -----
# Step 3: Insert sample data (10 rows)
# -----
sample_data = [
    ("Alice Johnson", "123 Maple St", "Springfield", "IL", "62704", "217-555-1234"),
    ("Bob Smith", "456 Oak St", "Madison", "WI", "53703", "608-555-5678"),
    ("Carla Lopez", "789 Pine Ave", "Denver", "CO", "80203", "303-555-9012"),
    ("David Lee", "101 Elm Dr", "Austin", "TX", "73301", "512-555-3456"),
    ("Emma Patel", "202 Birch Blvd", "Raleigh", "NC", "27601", "919-555-7890"),
    ("Frank Miller", "303 Cedar Ln", "Boise", "ID", "83702", "208-555-1122"),
    ("Grace Kim", "404 Aspen Way", "Seattle", "WA", "98101", "206-555-3344"),
    ("Henry Young", "505 Willow Rd", "Phoenix", "AZ", "85001", "602-555-5566"),
    ("Isabel Chen", "606 Sycamore Ct", "Boston", "MA", "02108", "617-555-7788"),
    ("Jake Nguyen", "707 Magnolia Pl", "Portland", "OR", "97201", "503-555-9900")
]

cursor.executemany("INSERT INTO contacts VALUES (?, ?, ?, ?, ?, ?)", sample_data)
conn.commit()

# -----
# Step 4: Query the table and display as DataFrame
# -----
df = pd.read_sql_query("SELECT * FROM contacts", conn)
print("\n All Contacts:")
print(df)
```

```
# Optional: Close the connection
conn.close()
```

All Contacts:

	name	address	city	state	zip	phone_number
0	Alice Johnson	123 Maple St	Springfield	IL	62704	217-555-1234
1	Bob Smith	456 Oak St	Madison	WI	53703	608-555-5678
2	Carla Lopez	789 Pine Ave	Denver	CO	80203	303-555-9012
3	David Lee	101 Elm Dr	Austin	TX	73301	512-555-3456
4	Emma Patel	202 Birch Blvd	Raleigh	NC	27601	919-555-7890
5	Frank Miller	303 Cedar Ln	Boise	ID	83702	208-555-1122
6	Grace Kim	404 Aspen Way	Seattle	WA	98101	206-555-3344
7	Henry Young	505 Willow Rd	Phoenix	AZ	85001	602-555-5566
8	Isabel Chen	606 Sycamore Ct	Boston	MA	02108	617-555-7788
9	Jake Nguyen	707 Magnolia Pl	Portland	OR	97201	503-555-9900
10	Alice Johnson	123 Maple St	Springfield	IL	62704	217-555-1234
11	Bob Smith	456 Oak St	Madison	WI	53703	608-555-5678
12	Carla Lopez	789 Pine Ave	Denver	CO	80203	303-555-9012
13	David Lee	101 Elm Dr	Austin	TX	73301	512-555-3456
14	Emma Patel	202 Birch Blvd	Raleigh	NC	27601	919-555-7890
15	Frank Miller	303 Cedar Ln	Boise	ID	83702	208-555-1122
16	Grace Kim	404 Aspen Way	Seattle	WA	98101	206-555-3344
17	Henry Young	505 Willow Rd	Phoenix	AZ	85001	602-555-5566
18	Isabel Chen	606 Sycamore Ct	Boston	MA	02108	617-555-7788
19	Jake Nguyen	707 Magnolia Pl	Portland	OR	97201	503-555-9900

In []: